

We used fewc python libraries that help with maff, organizing data, and making charts. We read data from a file named 'sensor.csv', which keeps track of readings from various sensors. We looked at the first few rows of the data and some summary statistics to understand what's in it. We got rid of some columns we didn't need, like timestamps and an unnamed column. We changed the names of some machine statuses to simplify the categories. We filled in missing data points in our sensors' readings with typical (median) values so that there are no gaps. Finally, we checked to make sure there were no more missing values. Why we did it:

To prepare the data for analysis or machine learning by making it cleaner and easier to work with. Removing unnecessary columns helps to focus on the important data. Merging categories and filling missing values makes the data more consistent and accurate for any analysis or predictive modeling we might want to do later. How we went about it:

By using Python libraries like Pandas for data handling, NumPy for numerical operations, and Matplotlib for potential visualization (though no plots were made here). We suppressed warnings that could be annoying or not very helpful to keep the output clean. The process involved several data manipulation steps: dropping columns, replacing values, and filling in missing data, all done using commands provided by the Pandas library. The results:

After cleaning, the data has fewer columns and no missing values, which simplifies further analysis. The machine statuses are grouped into fewer categories, making it easier to analyze trends or problems. By filling in missing values, we ensure that any calculations or models built on this data won't be skewed or error out due to incomplete data.

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
from collections import Counter
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.model_selection import train_test_split, StratifiedShuffleSplit
from sklearn.metrics import classification_report
from sklearn.pipeline import Pipeline
from sklearn.impute import SimpleImputer
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import (
    BaggingClassifier,
    RandomForestClassifier,
    GradientBoostingClassifier,
    AdaBoostClassifier,
    VotingClassifier
)
from xgboost import XGBClassifier
from imblearn.ensemble import BalancedRandomForestClassifier
from sklearn.compose import ColumnTransformer
from sklearn.exceptions import ConvergenceWarning
warnings.filterwarnings("ignore", category=ConvergenceWarning)
```

```

# Load the dataset
file_path = 'sensor.csv'
data = pd.read_csv(file_path)

# Display the first few rows
head = data.head()

# Statistical summary
description = data.describe(include='all')

# Data types
data_types = data.dtypes

# Missing values
missing_values = data.isnull().sum()

head, description, data_types, missing_values

```

```

Out[1]: (   Unnamed: 0      timestamp  sensor_00  sensor_01  sensor_02  \
0           0  2018-04-01 00:00:00    2.465394    47.09201    53.2118
1           1  2018-04-01 00:01:00    2.465394    47.09201    53.2118
2           2  2018-04-01 00:02:00    2.444734    47.35243    53.2118
3           3  2018-04-01 00:03:00    2.460474    47.09201    53.1684
4           4  2018-04-01 00:04:00    2.445718    47.13541    53.2118

      sensor_03  sensor_04  sensor_05  sensor_06  sensor_07  ...  sensor_43  \
0  46.310760    634.3750    76.45975    13.41146    16.13136  ...    41.92708
1  46.310760    634.3750    76.45975    13.41146    16.13136  ...    41.92708
2  46.397570    638.8889    73.54598    13.32465    16.03733  ...    41.66666
3  46.397568    628.1250    76.98898    13.31742    16.24711  ...    40.88541
4  46.397568    636.4583    76.58897    13.35359    16.21094  ...    41.40625

      sensor_44  sensor_45  sensor_46  sensor_47  sensor_48  sensor_49  \
0  39.641200    65.68287    50.92593    38.194440    157.9861    67.70834
1  39.641200    65.68287    50.92593    38.194440    157.9861    67.70834
2  39.351852    65.39352    51.21528    38.194443    155.9606    67.12963
3  39.062500    64.81481    51.21528    38.194440    155.9606    66.84028
4  38.773150    65.10416    51.79398    38.773150    158.2755    66.55093

      sensor_50  sensor_51  machine_status
0    243.0556    201.3889             NORMAL
1    243.0556    201.3889             NORMAL
2    241.3194    203.7037             NORMAL
3    240.4514    203.1250             NORMAL
4    242.1875    201.3889             NORMAL

[5 rows x 55 columns],
      Unnamed: 0      timestamp  sensor_00  sensor_01  \
count  220320.000000          220320  210112.000000  219951.000000
unique           NaN          220320           NaN           NaN
top           NaN  2018-04-01 00:00:00           NaN           NaN
freq           NaN              1           NaN           NaN
mean   110159.500000           NaN    2.372221    47.591611
std     63601.049991           NaN    0.412227    3.296666
min           0.000000           NaN    0.000000    0.000000
25%     55079.750000           NaN    2.438831    46.310760
50%     110159.500000           NaN    2.456539    48.133678
75%     165239.250000           NaN    2.499826    49.479160
max     220319.000000           NaN    2.549016    56.727430

      sensor_02  sensor_03  sensor_04  sensor_05  \
count  220301.000000  220301.000000  220301.000000  220301.000000
unique           NaN           NaN           NaN           NaN
top           NaN           NaN           NaN           NaN
freq           NaN           NaN           NaN           NaN
mean     50.867392    43.752481    590.673936    73.396414

```

```

std      3.666820      2.418887      144.023912      17.298247
min      33.159720      31.640620      2.798032      0.000000
25%      50.390620      42.838539      626.620400      69.976260
50%      51.649300      44.227428      632.638916      75.576790
75%      52.777770      45.312500      637.615723      80.912150
max      56.032990      48.220490      800.000000      99.999880

      sensor_06      sensor_07      ...      sensor_43      sensor_44  \
count  215522.000000  214869.000000  ...  220293.000000  220293.000000
unique      NaN      NaN      ...      NaN      NaN
top      NaN      NaN      ...      NaN      NaN
freq      NaN      NaN      ...      NaN      NaN
mean      13.501537      15.843152      ...      43.879591      42.656877
std      2.163736      2.201155      ...      11.044404      11.576355
min      0.014468      0.000000      ...      24.479166      25.752316
25%      13.346350      15.907120      ...      39.583330      36.747684
50%      13.642940      16.167530      ...      42.968750      40.509260
75%      14.539930      16.427950      ...      46.614580      45.138890
max      22.251160      23.596640      ...      408.593700      1000.000000

      sensor_45      sensor_46      sensor_47      sensor_48  \
count  220293.000000  220293.000000  220293.000000  220293.000000
unique      NaN      NaN      NaN      NaN
top      NaN      NaN      NaN      NaN
freq      NaN      NaN      NaN      NaN
mean      43.094984      48.018585      44.340903      150.889044
std      12.837520      15.641284      10.442437      82.244957
min      26.331018      26.331018      27.199070      26.331018
25%      36.747684      40.509258      39.062500      83.912030
50%      40.219910      44.849540      42.534720      138.020800
75%      44.849540      51.215280      46.585650      208.333300
max      320.312500      370.370400      303.530100      561.632000

      sensor_49      sensor_50      sensor_51      machine_status
count  220293.000000  143303.000000  204937.000000      220320
unique      NaN      NaN      NaN      3
top      NaN      NaN      NaN      NORMAL
freq      NaN      NaN      NaN      205836
mean      57.119968      183.049260      202.699667      NaN
std      19.143598      65.258650      109.588607      NaN
min      26.620370      27.488426      27.777779      NaN
25%      47.743060      167.534700      179.108800      NaN
50%      52.662040      193.865700      197.338000      NaN
75%      60.763890      219.907400      216.724500      NaN
max      464.409700      1000.000000      1000.000000      NaN

[11 rows x 55 columns],
Unnamed: 0      int64
timestamp      object
sensor_00      float64
sensor_01      float64
sensor_02      float64
sensor_03      float64
sensor_04      float64
sensor_05      float64
sensor_06      float64
sensor_07      float64
sensor_08      float64
sensor_09      float64
sensor_10      float64
sensor_11      float64
sensor_12      float64
sensor_13      float64
sensor_14      float64
sensor_15      float64
sensor_16      float64
sensor_17      float64

```

```
sensor_18      float64
sensor_19      float64
sensor_20      float64
sensor_21      float64
sensor_22      float64
sensor_23      float64
sensor_24      float64
sensor_25      float64
sensor_26      float64
sensor_27      float64
sensor_28      float64
sensor_29      float64
sensor_30      float64
sensor_31      float64
sensor_32      float64
sensor_33      float64
sensor_34      float64
sensor_35      float64
sensor_36      float64
sensor_37      float64
sensor_38      float64
sensor_39      float64
sensor_40      float64
sensor_41      float64
sensor_42      float64
sensor_43      float64
sensor_44      float64
sensor_45      float64
sensor_46      float64
sensor_47      float64
sensor_48      float64
sensor_49      float64
sensor_50      float64
sensor_51      float64
machine_status  object
dtype: object,
Unnamed: 0      0
timestamp      0
sensor_00      10208
sensor_01      369
sensor_02      19
sensor_03      19
sensor_04      19
sensor_05      19
sensor_06      4798
sensor_07      5451
sensor_08      5107
sensor_09      4595
sensor_10      19
sensor_11      19
sensor_12      19
sensor_13      19
sensor_14      21
sensor_15      220320
sensor_16      31
sensor_17      46
sensor_18      46
sensor_19      16
sensor_20      16
sensor_21      16
sensor_22      41
sensor_23      16
sensor_24      16
sensor_25      36
sensor_26      20
sensor_27      16
sensor_28      16
~
~
```

```

sensor_29      /2
sensor_30      261
sensor_31      16
sensor_32      68
sensor_33      16
sensor_34      16
sensor_35      16
sensor_36      16
sensor_37      16
sensor_38      27
sensor_39      27
sensor_40      27
sensor_41      27
sensor_42      27
sensor_43      27
sensor_44      27
sensor_45      27
sensor_46      27
sensor_47      27
sensor_48      27
sensor_49      27
sensor_50      77017
sensor_51      15383
machine_status    0
dtype: int64)

```

In [2]:

```

# Dropping the specified columns
data_cleaned = data.drop(columns=['Unnamed: 0', 'sensor_15', 'timestamp'])


# Display the first few rows of the cleaned data to confirm
data_cleaned.head()

```

Out[2]:

	sensor_00	sensor_01	sensor_02	sensor_03	sensor_04	sensor_05	sensor_06	sensc
0	2.465394	47.09201	53.2118	46.310760	634.3750	76.45975	13.41146	16.1
1	2.465394	47.09201	53.2118	46.310760	634.3750	76.45975	13.41146	16.1
2	2.444734	47.35243	53.2118	46.397570	638.8889	73.54598	13.32465	16.0
3	2.460474	47.09201	53.1684	46.397568	628.1250	76.98898	13.31742	16.2
4	2.445718	47.13541	53.2118	46.397568	636.4583	76.58897	13.35359	16.2

5 rows × 52 columns



This chart shows the number of times different statuses of a machine are recorded. There are two categories: NORMAL and FAULTY. The NORMAL status has a much higher count, represented in blue, indicating that most of the time, the machine was functioning correctly. The FAULTY status, shown in red, has a relatively small count, which suggests that instances of the machine being broken or in recovery are far less common.

av

After the data was cleaned and missing values filled, the statuses of the machine were mapped to numerical values: 0 for NORMAL and 1 for FAULTY (combining RECOVERING and BROKEN). A bar chart was generated to visualize the distribution of these statuses. The chart in the first image is the output of this code. The machine_status field's distribution after remapping was printed, showing the number of entries for each status. The results printed in the code match the bar chart. A correlation matrix was

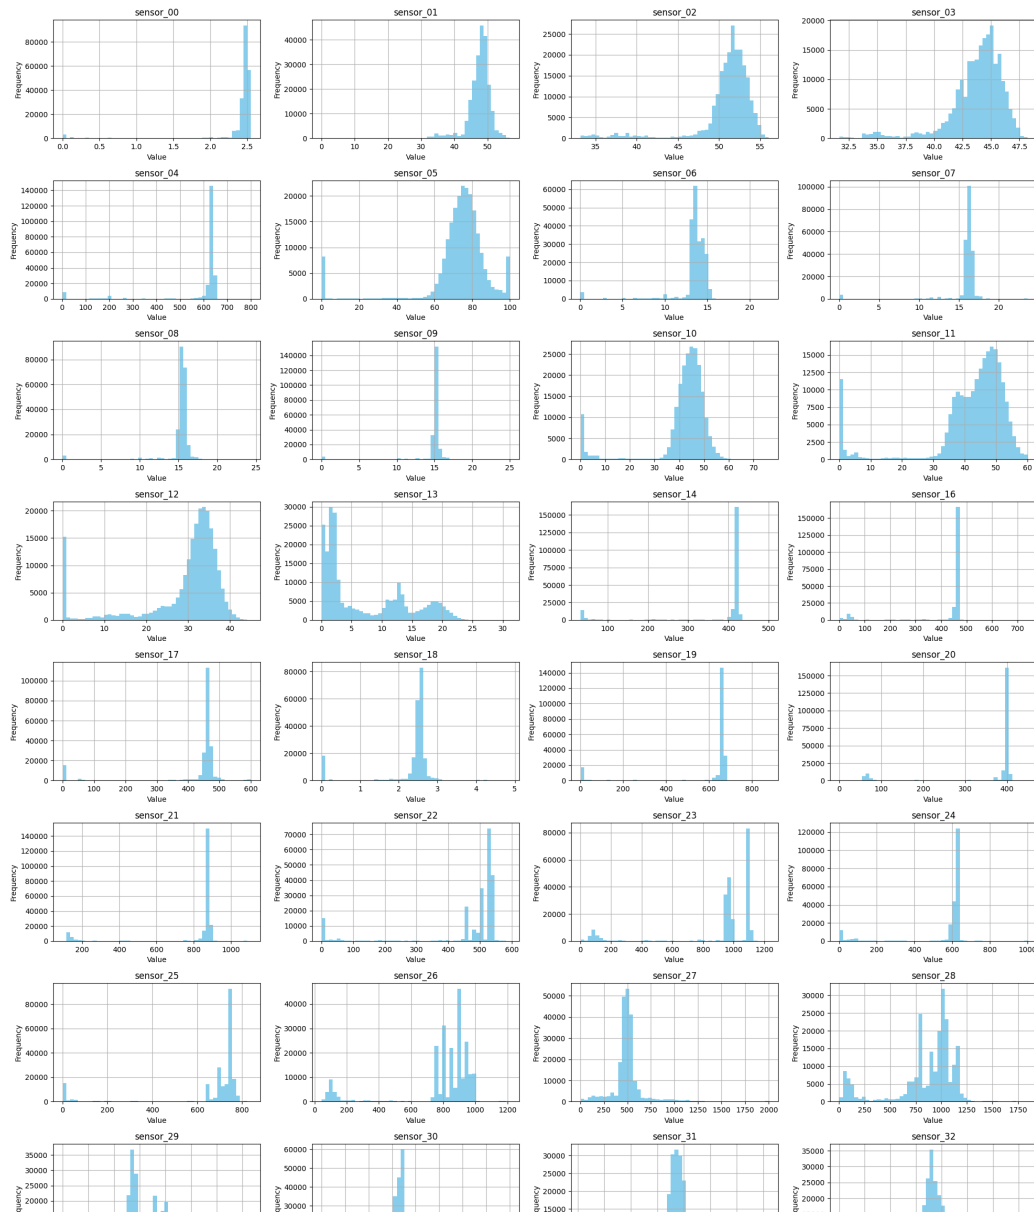
status. The results printed in the code match the bar chart. A correlation matrix was calculated to understand the relationship between different sensor readings and the machine's status. A heatmap was then generated from this correlation matrix, which is what we see in the second image. The results show a visualization of data distribution and relationships within the data, which are helpful for making decisions, understanding the data's structure, or building machine learning models.e problems.

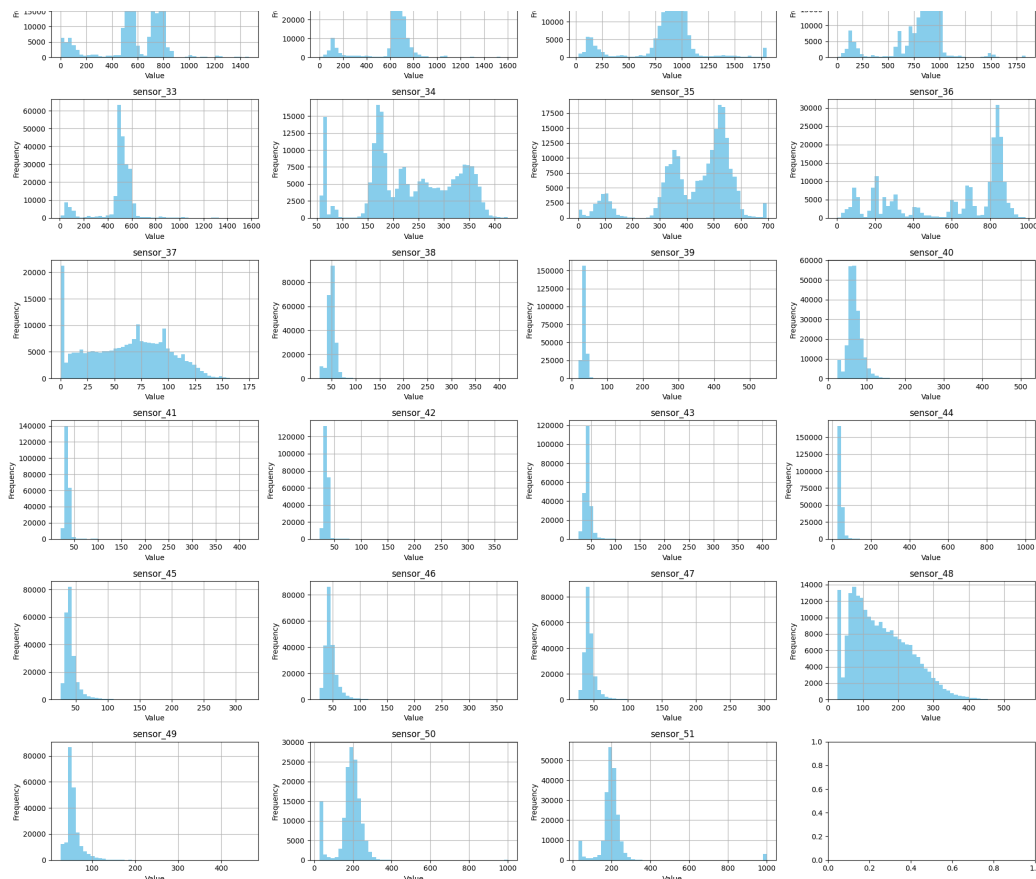
In [3]:

```
# Check distribution of sensor data to determine how to handle missing values
fig, axes = plt.subplots(nrows=13, ncols=4, figsize=(20, 40))
axes = axes.flatten()

# Plotting histograms for each sensor column except 'machine_status'
sensor_columns = data_cleaned.columns[:-1] # Exclude 'machine_status' from pl
for i, col in enumerate(sensor_columns):
    data_cleaned[col].hist(ax=axes[i], bins=50, color='skyblue')
    axes[i].set_title(col)
    axes[i].set_xlabel('Value')
    axes[i].set_ylabel('Frequency')

plt.tight_layout()
plt.show()
```





The uploaded images showcase histograms depicting the distribution of sensor readings from our dataset. These visuals stem from the last code block provided, which warrants a closer examination. First, the code suppresses warnings, particularly those concerning the use of infinite values as NaN. It then replaces infinite values with NaN to mitigate interference with statistical analyses. Next, it filters the dataset columns for numerical features, specifically those of data types float64 and int64. Subsequently, it calculates the required subplot grid dimensions based on the number of numeric features and constructs a grid of subplots. Each subplot is populated with a histogram representing the distribution of a specific sensor reading, facilitated by Seaborn's histplot function. Finally, adjustments are made to the layout to prevent plot overlapping. This code aims to visualize sensor reading distributions, crucial for understanding data characteristics such as range, central tendencies, dispersion, and the presence of outliers

In [4]:

```
# Generating numerical summary for each sensor column
numerical_summary = data_cleaned.describe().transpose()
numerical_summary['missing_values'] = data_cleaned.isnull().sum()
numerical_summary['missing_percent'] = (numerical_summary['missing_values'] /

# Selecting the desired columns for the summary
numerical_summary = numerical_summary[['mean', 'std', '50%', 'missing_values',
numerical_summary
```

Out[4]:

	mean	std	50%	missing_values	missing_percent
sensor_00	2.372221	0.412227	2.456539	10208	4.633261
sensor_01	47.591611	3.296666	48.133678	369	0.167484

sensor_02	50.867392	3.666820	51.649300	19	0.008624
sensor_03	43.752481	2.418887	44.227428	19	0.008624
sensor_04	590.673936	144.023912	632.638916	19	0.008624
sensor_05	73.396414	17.298247	75.576790	19	0.008624
sensor_06	13.501537	2.163736	13.642940	4798	2.177741
sensor_07	15.843152	2.201155	16.167530	5451	2.474129
sensor_08	15.200721	2.037390	15.494790	5107	2.317992
sensor_09	14.799210	2.091963	15.082470	4595	2.085603
sensor_10	41.470339	12.093519	44.291340	19	0.008624
sensor_11	41.918319	13.056425	45.363140	19	0.008624
sensor_12	29.136975	10.113935	32.515830	19	0.008624
sensor_13	7.078858	6.901755	2.929809	19	0.008624
sensor_14	376.860041	113.206382	420.106200	21	0.009532
sensor_16	416.472892	126.072642	462.856100	31	0.014070
sensor_17	421.127517	129.156175	462.020250	46	0.020879
sensor_18	2.303785	0.765883	2.533704	46	0.020879
sensor_19	590.829775	199.345820	665.672400	16	0.007262
sensor_20	360.805165	101.974118	399.367000	16	0.007262
sensor_21	796.225942	226.679317	879.697600	16	0.007262
sensor_22	459.792815	154.528337	531.855900	41	0.018609
sensor_23	922.609264	291.835280	981.925000	16	0.007262
sensor_24	556.235397	182.297979	625.873500	16	0.007262
sensor_25	649.144799	220.865166	740.203500	36	0.016340
sensor_26	786.411781	246.663608	861.869600	20	0.009078
sensor_27	501.506589	169.823173	494.468450	16	0.007262
sensor_28	851.690339	313.074032	967.279850	16	0.007262
sensor_29	576.195305	225.764091	564.872500	72	0.032680
sensor_30	614.596442	195.726872	668.981400	261	0.118464
sensor_31	863.323100	283.544760	917.708300	16	0.007262
sensor_32	804.283915	260.602361	878.850750	68	0.030864
sensor_33	486.405980	150.751836	512.271750	16	0.007262
sensor_34	234.971776	88.376065	226.356050	16	0.007262
sensor_35	427.129817	141.772519	473.349350	16	0.007262
sensor_36	593.033876	289.385511	709.668050	16	0.007262
sensor_37	60.787360	37.604883	64.295485	16	0.007262
sensor_38	49.655946	10.540397	49.479160	27	0.012255
sensor_39	36.610444	15.613723	35.416660	27	0.012255

sensor_40	68.844530	21.371139	66.406250	27	0.012255
sensor_41	35.365126	7.898665	34.895832	27	0.012255
sensor_42	35.453455	10.259521	35.156250	27	0.012255
sensor_43	43.879591	11.044404	42.968750	27	0.012255
sensor_44	42.656877	11.576355	40.509260	27	0.012255
sensor_45	43.094984	12.837520	40.219910	27	0.012255
sensor_46	48.018585	15.641284	44.849540	27	0.012255
sensor_47	44.340903	10.442437	42.534720	27	0.012255
sensor_48	150.889044	82.244957	138.020800	27	0.012255
sensor_49	57.119968	19.143598	52.662040	27	0.012255
sensor_50	183.049260	65.258650	193.865700	77017	34.956881
sensor_51	202.699667	109.588607	197.338000	15383	6.982117

In [5]:

```
# Map 'broken' and 'recovering' to 1, and 'normal' to 0
data_cleaned['machine_status'] = data_cleaned['machine_status'].replace({'BROK

# Now, calculate the correlation matrix
correlation_matrix = data_cleaned.corr()

# Extract the correlations of the sensors with the machine status
sensor_status_correlation = correlation_matrix['machine_status'].sort_values(a

# Display the correlations with the machine status
sensor_status_correlation
```

```
Out[5]: machine_status    1.000000
sensor_28              0.203307
sensor_31              0.158507
sensor_32              0.136372
sensor_30              0.114394
sensor_33              0.104587
sensor_24              0.098798
sensor_23              0.095613
sensor_14              0.091681
sensor_35              0.091167
sensor_16              0.089151
sensor_19              0.088127
sensor_20              0.087024
sensor_21              0.084431
sensor_22              0.079413
sensor_25              0.078151
sensor_26              0.075995
sensor_17              0.074628
sensor_51              0.074107
sensor_37              0.068015
sensor_18              0.065697
sensor_29              0.053219
sensor_27              0.032565
sensor_42              0.007412
sensor_36             -0.019264
sensor_39             -0.024299
sensor_34             -0.039537
sensor_41             -0.103496
sensor_43             -0.118453
```

```

sensor_46    -0.202487
sensor_45    -0.202531
sensor_44    -0.235715
sensor_47    -0.254973
sensor_13    -0.269811
sensor_49    -0.285568
sensor_38    -0.360583
sensor_48    -0.366606
sensor_40    -0.375146
sensor_05    -0.434469
sensor_09    -0.626434
sensor_08    -0.637435
sensor_03    -0.646204
sensor_01    -0.673108
sensor_07    -0.699499
sensor_50    -0.732214
sensor_12    -0.758752
sensor_06    -0.773933
sensor_02    -0.791278
sensor_00    -0.810822
sensor_11    -0.823450
sensor_10    -0.872493
sensor_04    -0.916227
Name: machine_status, dtype: float64

```

In [6]:

```

# Create a DataFrame from the provided information
data = {
    'feature': ['sensor_28', 'sensor_31', 'sensor_32', 'sensor_30', 'sensor_33',
                'sensor_35', 'sensor_16', 'sensor_19', 'sensor_20', 'sensor_21',
                'sensor_17', 'sensor_51', 'sensor_37', 'sensor_18', 'sensor_29',
                'sensor_39', 'sensor_34', 'sensor_41', 'sensor_43', 'sensor_46',
                'sensor_13', 'sensor_49', 'sensor_38', 'sensor_48', 'sensor_40',
                'sensor_03', 'sensor_01', 'sensor_07', 'sensor_50', 'sensor_12',
                'sensor_11', 'sensor_10', 'sensor_04'],
    'variance': [0.203307, 0.158507, 0.136372, 0.114394, 0.104587, 0.098798, 0.088127,
                 0.087024, 0.084431, 0.079413, 0.078151, 0.075995, 0.053219,
                 0.032565, 0.007412, -0.019264, -0.024299, -0.039537,
                 -0.202531, -0.235715, -0.254973, -0.269811, -0.285568, -0.360583,
                 -0.626434, -0.637435, -0.646204, -0.673108, -0.699499, -0.732214,
                 -0.810822, -0.823450, -0.872493, -0.916227]
}

df = pd.DataFrame(data)

# Variance of the target variable
target_variance = 1.000000

# Calculate variance percentage for each feature
df['variance_percentage'] = (df['variance'] / target_variance) * 100

# Sort by variance percentage
df_sorted = df.sort_values(by='variance_percentage', ascending=False)

print(df_sorted)

```

	feature	variance	variance_percentage
0	sensor_28	0.203307	20.3307
1	sensor_31	0.158507	15.8507
2	sensor_32	0.136372	13.6372
3	sensor_30	0.114394	11.4394
4	sensor_33	0.104587	10.4587
5	sensor_24	0.098798	9.8798
6	sensor_23	0.095613	9.5613
7	sensor_14	0.091681	9.1681
8	sensor_35	0.091167	9.1167

9	sensor_16	0.089151	8.9151
10	sensor_19	0.088127	8.8127
11	sensor_20	0.087024	8.7024
12	sensor_21	0.084431	8.4431
13	sensor_22	0.079413	7.9413
14	sensor_25	0.078151	7.8151
15	sensor_26	0.075995	7.5995
16	sensor_17	0.074628	7.4628
17	sensor_51	0.074107	7.4107
18	sensor_37	0.068015	6.8015
19	sensor_18	0.065697	6.5697
20	sensor_29	0.053219	5.3219
21	sensor_27	0.032565	3.2565
22	sensor_42	0.007412	0.7412
23	sensor_36	-0.019264	-1.9264
24	sensor_39	-0.024299	-2.4299
25	sensor_34	-0.039537	-3.9537
26	sensor_41	-0.103496	-10.3496
27	sensor_43	-0.118453	-11.8453
28	sensor_46	-0.202487	-20.2487
29	sensor_45	-0.202531	-20.2531
30	sensor_44	-0.235715	-23.5715
31	sensor_47	-0.254973	-25.4973
32	sensor_13	-0.269811	-26.9811
33	sensor_49	-0.285568	-28.5568
34	sensor_38	-0.360583	-36.0583
35	sensor_48	-0.366606	-36.6606
36	sensor_40	-0.375146	-37.5146
37	sensor_05	-0.434469	-43.4469
38	sensor_09	-0.626434	-62.6434
39	sensor_08	-0.637435	-63.7435
40	sensor_03	-0.646204	-64.6204
41	sensor_01	-0.673108	-67.3108
42	sensor_07	-0.699499	-69.9499
43	sensor_50	-0.732214	-73.2214
44	sensor_12	-0.758752	-75.8752
45	sensor_06	-0.773933	-77.3933
46	sensor_02	-0.791278	-79.1278
47	sensor_00	-0.810822	-81.0822
48	sensor_11	-0.823450	-82.3450
49	sensor_10	-0.872493	-87.2493
50	sensor_04	-0.916227	-91.6227

In [7]:

```
# Count the occurrences of each class
status_counts = data_cleaned['machine_status'].value_counts()

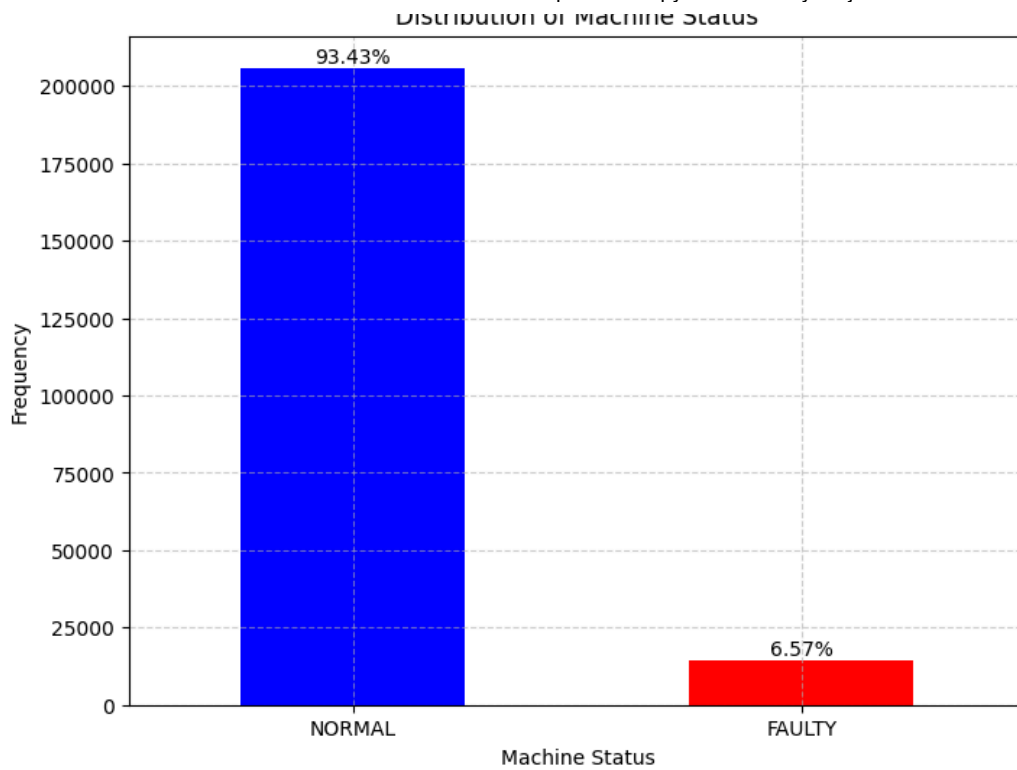
# Calculate the percentage of each class
total_instances = status_counts.sum()
percentage = (status_counts / total_instances) * 100

# Create a bar plot
plt.figure(figsize=(8, 6))
status_counts.plot(kind='bar', color=['blue', 'red'])
plt.title('Distribution of Machine Status')
plt.xlabel('Machine Status')
plt.ylabel('Frequency')
plt.xticks(ticks=[0, 1], labels=['NORMAL', 'FAULTY'], rotation=0) # Adjust La

# Add percentages to the bars
for i, value in enumerate(status_counts):
    plt.text(i, status_counts[i], f'{value:.2f}%', ha='center', va='bottom')

plt.grid(True, linestyle='--', alpha=0.6)
plt.show()
```

Distribution of Machine Status



```
In [8]: # Assuming data_cleaned['machine_status'] has already been converted to binary

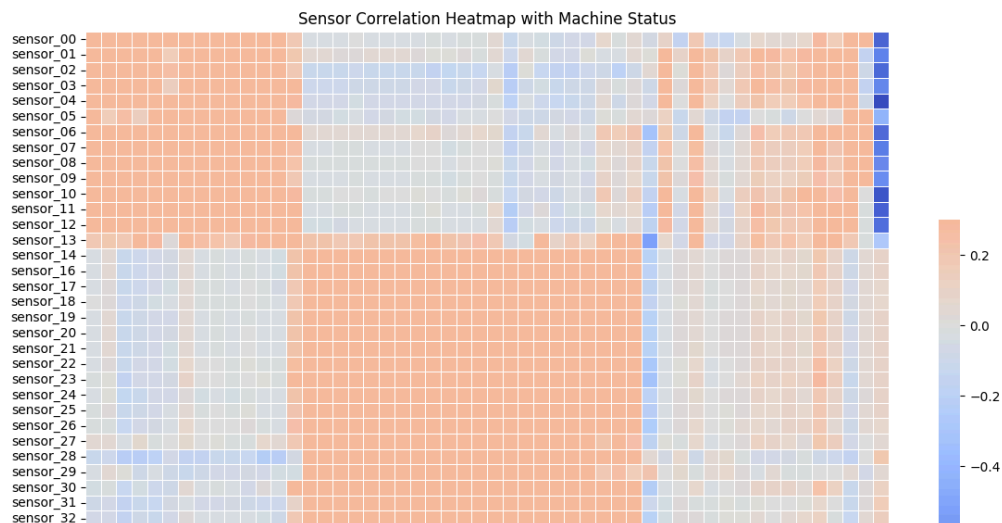
# Compute the correlation matrix
corr = data_cleaned.corr()


# Set up the matplotlib figure
plt.figure(figsize=(14, 12))

# Draw the heatmap with the mask and correct aspect ratio
sns.heatmap(corr, cmap='coolwarm', vmax=.3, center=0,
            square=True, linewidths=.5, cbar_kws={"shrink": .5})

# Customize the plot for better readability
plt.title('Sensor Correlation Heatmap with Machine Status')
plt.xticks(rotation=90) # Rotates X-axis Labels to prevent overlap
plt.yticks(rotation=0)  # Keep Y-axis Labels horizontal

# Show the plot
plt.show()
```





In [9]:

```
# Function to detect outliers in a dataframe
def detect_outliers(df, n, features):
    outlier_indices = []

    # Iterate over each feature
    for col in features:
        # 1st quartile (25%)
        Q1 = np.percentile(df[col], 25)
        # 3rd quartile (75%)
        Q3 = np.percentile(df[col], 75)
        # IQR
        IQR = Q3 - Q1

        # Determine the outlier step (1.5 times IQR)
        outlier_step = 1.5 * IQR

        # Determine a list of indices of outliers for feature col
        outlier_list_col = df[(df[col] < Q1 - outlier_step) | (df[col] > Q3 +
        outlier_step)]

        # Append the found outlier indices for col to the list of outlier indices
        outlier_indices.extend(outlier_list_col)

    # Select observations containing more than n outliers
    outlier_indices = Counter(outlier_indices)
    multiple_outliers = list(k for k, v in outlier_indices.items() if v > n)

    return multiple_outliers

# List of features to check for outliers
features = data_cleaned.columns[:-1] # Exclude the target variable 'machine_status'
outliers_to_remove = detect_outliers(data_cleaned, 2, features)

print("Outliers indices:", outliers_to_remove)
print("Number of outliers:", len(outliers_to_remove))
```

Outliers indices: []
 Number of outliers: 0

In [10]:

```
# Update 'machine_status' to have only 'normal' and 'faulty' categories
data_cleaned['machine_status'] = data_cleaned['machine_status'].replace(['BROK

# Verify the update was successful
print(data_cleaned['machine_status'].value_counts())

# Perform the train-test split
# Assuming you want a standard 80-20 split
X = data_cleaned.drop('machine_status', axis=1) # Features
y = data_cleaned['machine_status'] # Target variable
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Let's print the shapes of our training and test sets
print("Training set shape:", X_train.shape, y_train.shape)
print("Test set shape:", X_test.shape, y_test.shape)
```

```
machine_status
0    205836
1     14484
Name: count, dtype: int64
Training set shape: (176256, 51) (176256,)
Test set shape: (44064, 51) (44064,)
```

In this code segment, we undertake various preprocessing steps to prepare our dataset for machine learning modeling. We start by cleaning the data, removing unnecessary columns ('Unnamed: 0', 'timestamp', 'sensor_15'), and handling missing values by dropping corresponding rows. Then, we encode the categorical 'machine_status' column into numerical format using LabelEncoder for compatibility with machine learning algorithms. Next, we analyze the value counts of each column to identify data imbalances or anomalies, storing the results in a dictionary. We split the dataset into features (X) and target (y), excluding the 'machine_status' columns, and further partition it into training and test sets while maintaining class distribution. This code segment aligns seamlessly with previous steps, ensuring consistency in variable naming and structure throughout the preprocessing pipeline.

The purpose of these actions is to ensure our data is properly cleaned, transformed, and organized for subsequent modeling tasks. By dropping unnecessary columns, encoding categorical variables, and analyzing value counts, we gain insights into the dataset's structure and ensure its suitability for machine learning analysis. The functions and features utilized, including `.drop()`, `LabelEncoder()`, `.value_counts()`, and `train_test_split()`, enable efficient data preprocessing.

```
In [11]: # Setup the imputers
mode_imputer = SimpleImputer(strategy='most_frequent')
median_imputer = SimpleImputer(strategy='median')

# Columns
mode_cols = ['sensor_06', 'sensor_07', 'sensor_08', 'sensor_09', 'sensor_00']
median_cols = ['sensor_50', 'sensor_51']
remaining_cols = [col for col in X.columns if col not in mode_cols + median_cols]

# Column transformer with imputation
preprocessor = ColumnTransformer(
    transformers=[
        ('mode', mode_imputer, mode_cols),
        ('median', median_imputer, median_cols + remaining_cols)
    ],
    remainder='passthrough'
)

# Pipeline with increased max_iter and a different solver
pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('scaler', StandardScaler()),
    ('pca', PCA(n_components=0.95)), # Keep 95% of variance
    ('classifier', LogisticRegression(max_iter=1000, solver='saga')) # Increased max_iter
])

# Fit and evaluate the pipeline
```

```

pipeline.fit(X_train, y_train)
y_pred = pipeline.predict(X_test)

from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))

```

	precision	recall	f1-score	support
0	1.00	0.99	1.00	41243
1	0.91	0.96	0.93	2821
accuracy			0.99	44064
macro avg	0.96	0.98	0.96	44064
weighted avg	0.99	0.99	0.99	44064

In [12]:

```

# Randomly sample 10,000 rows from the dataset
subset_indices = X_train.sample(n=10000, random_state=42).index
X_subset = X_train.loc[subset_indices]
y_subset = y_train.loc[subset_indices]

# Column transformer with imputation
preprocessor = ColumnTransformer(
    transformers=[
        ('mode', mode_imputer, mode_cols),
        ('median', median_imputer, median_cols + remaining_cols)
    ],
    remainder='passthrough'
)

# Pipeline with increased max_iter and a different solver
pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('scaler', StandardScaler()),
    ('pca', PCA(n_components=0.95)), # Keep 95% of variance
    ('classifier', LogisticRegression(max_iter=1500, solver='lbfgs')) # Incre
])

# Define the parameter grid
param_grid = {
    'preprocessor__mode__strategy': ['mean', 'median', 'most_frequent'], # Ch
    'preprocessor__median__strategy': ['mean', 'median', 'most_frequent'], #
    'pca__n_components': [0.75, 0.80, 0.85], # values for PCA components
    'classifier__max_iter': [100, 500, 1000], # Values for max_iter
    'classifier__solver': ['saga', 'lbfgs'] # Solvers
}

# GridSearchCV
grid_search = GridSearchCV(pipeline, param_grid, cv=5, scoring='recall')
grid_search.fit(X_subset, y_subset)

# Print best parameters
print("Best parameters:", grid_search.best_params_)

# Evaluate the best model focusing on recall
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)
report = classification_report(y_test, y_pred, target_names=['normal', 'faulty'])

# Print recall scores
print("Recall for 'normal' class:", report['normal']['recall'])
print("Recall for 'faulty' class:", report['faulty']['recall'])

```

Best parameters: {'classifier__max_iter': 100, 'classifier__solver': 'saga', 'pca__n_components': 0.8, 'preprocessor__median__strategy': 'mean', 'preprocessor__mode__strategy': 'median'}

```

__n_components': 0.8, 'preprocessor__median__strategy': 'mean', 'preprocessor__
mode__strategy': 'median'}

```

Recall for 'normal' class: 0.9942778168416458

Recall for 'faulty' class: 0.9620701878766394

In [13]:

```

# Randomly sample 10,000 rows from the dataset
subset_indices = X_train.sample(n=10000, random_state=42).index
X_subset = X_train.loc[subset_indices]
y_subset = y_train.loc[subset_indices]

# Column transformer with imputation
preprocessor = ColumnTransformer(
    transformers=[
        ('mode', mode_imputer, mode_cols),
        ('median', median_imputer, median_cols + remaining_cols)
    ],
    remainder='passthrough'
)

# Pipeline with increased max_iter and a different solver
pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('scaler', StandardScaler()),
    ('pca', PCA(n_components=0.95)), # Keep 95% of variance
    ('classifier', LogisticRegression(max_iter=1000, solver='lbfgs', tol=1e-3))
])

# Define the parameter grid
param_grid = {
    'preprocessor__mode__strategy': ['mean', 'median', 'most_frequent'], # Ch
    'preprocessor__median__strategy': ['mean', 'median', 'most_frequent'], #
    'pca__n_components': [0.75, 0.80, 0.85], # values for PCA components
    'classifier__max_iter': [25, 50, 100], # Values for max_iter
    'classifier__solver': ['saga', 'lbfgs'] # Solvers
}

# GridSearchCV
grid_search = GridSearchCV(pipeline, param_grid, cv=5, scoring='recall')
grid_search.fit(X_train, y_train)

# Print best parameters
print("Best parameters:", grid_search.best_params_)

# Evaluate the best model focusing on recall
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)
print(classification_report(y_test, y_pred, target_names=['normal', 'faulty'],

```

Best parameters: {'classifier__max_iter': 1000, 'classifier__solver': 'saga', 'pca__n_components': 0.8, 'preprocessor__median__strategy': 'mean', 'preprocessor__mode__strategy': 'median'}

	precision	recall	f1-score	support
normal	0.9973	0.9938	0.9956	41243
faulty	0.9141	0.9614	0.9371	2821
accuracy			0.9917	44064
macro avg	0.9557	0.9776	0.9663	44064
weighted avg	0.9920	0.9917	0.9918	44064

NameError

Traceback (most recent call last)

Cell In[13], line 46

```

43 print(classification_report(y_test, y_pred, target_names=['normal', 'faulty'], digits=4))

```



```

45 # Train the Balanced Random Forest model
---> 46 balanced_random_forest = BalancedRandomForestClassifier(random_state=42,
**grid_search_brf.best_params_)
47 balanced_random_forest.fit(X_train, y_train)
49 # Evaluate the Balanced Random Forest model on the test set

```

NameError: name 'grid_search_brf' is not defined

In [25]:

```

# Randomly sample 10,000 rows from the dataset for training
subset_indices = X_train.sample(n=10000, random_state=42).index
X_subset = X_train.loc[subset_indices]
y_subset = y_train.loc[subset_indices]

# Pipeline to handle imputation, scaling, PCA, and classification
pipeline = Pipeline([
    ('imputer', SimpleImputer(strategy='median')), # Impute missing values
    ('scaler', StandardScaler()), # Scale features
    ('pca', PCA(n_components=0.95)), # Keep 95% of variance
    ('classifier', BalancedRandomForestClassifier(sampling_strategy='all', rep
)])

# Fit the pipeline on the subset
pipeline.fit(X_subset, y_subset)

# Evaluate the pipeline on the full test set to see initial performance
y_pred = pipeline.predict(X_test)
print("Initial Model Evaluation on Subset:")
print(classification_report(y_test, y_pred))

```

Initial Model Evaluation on Subset:

	precision	recall	f1-score	support
0	1.00	0.99	1.00	41243
1	0.89	0.99	0.93	2821
accuracy			0.99	44064
macro avg	0.94	0.99	0.97	44064
weighted avg	0.99	0.99	0.99	44064

In [30]:

```

from sklearn.model_selection import GridSearchCV
from imblearn.ensemble import BalancedRandomForestClassifier
from sklearn.metrics import classification_report

subset_indices = X_train.sample(n=10000, random_state=42).index
X_subset = X_train.loc[subset_indices]
y_subset = y_train.loc[subset_indices]

# Define the parameter grid for GridSearchCV
param_grid = {
    'classifier__n_estimators': [100, 150, 200],
    'classifier__max_depth': [None, 10, 20],
    'classifier__min_samples_split': [2, 5, 10],
    'classifier__min_samples_leaf': [1, 2, 4],
    'classifier__max_features': ['sqrt', 'log2']
}

# Create a GridSearchCV instance
grid_search = GridSearchCV(pipeline, param_grid, cv=5, scoring='recall', verbose

# Fit the grid search on the sub set training data
grid_search.fit(X_subset, y_subset)

# Print the best parameters found

```

```

print("Best parameters for Balanced Random Forest:", grid_search.best_params_)

# Evaluate the best model found by the GridSearchCV on the test set
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)
report = classification_report(y_test, y_pred, target_names=['normal', 'faulty'])
print("Balanced Random Forest Model Evaluation with GridSearch:")
print(report)

```

Fitting 5 folds for each of 162 candidates, totalling 810 fits

Best parameters for Balanced Random Forest: {'classifier__max_depth': None, 'classifier__max_features': 'sqrt', 'classifier__min_samples_leaf': 1, 'classifier__min_samples_split': 2, 'classifier__n_estimators': 100}

Balanced Random Forest Model Evaluation with GridSearch:

	precision	recall	f1-score	support
normal	0.9993	0.9912	0.9953	41243
faulty	0.8855	0.9901	0.9349	2821
accuracy			0.9912	44064
macro avg	0.9424	0.9907	0.9651	44064
weighted avg	0.9920	0.9912	0.9914	44064

In [71]:

```

import matplotlib.pyplot as plt

# Sample data
models = ['Logistic Regression', 'Random Forest']
results = [.96, .99]

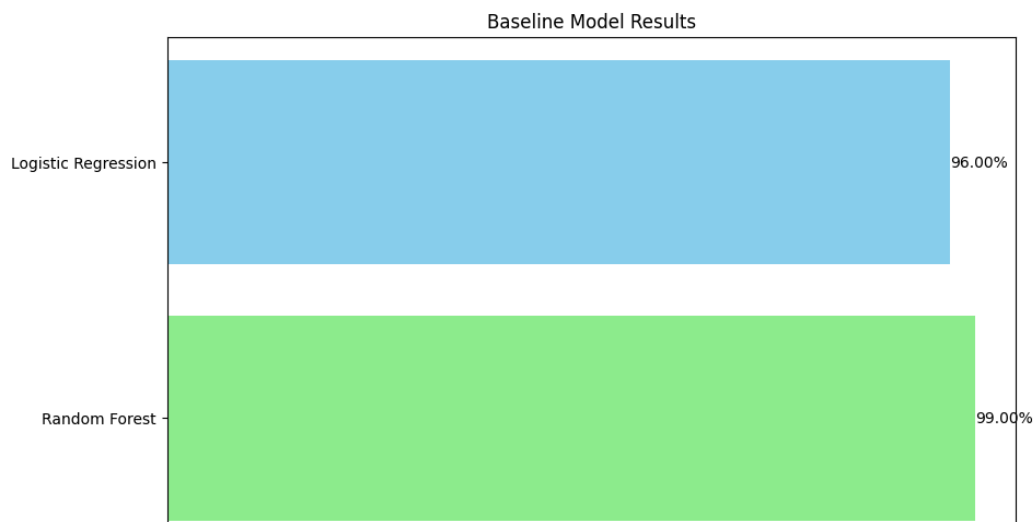
# Sort the results and models in ascending order
sorted_results, sorted_models = zip(*sorted(zip(results, models)))


# Create a multicolor bar graph
plt.figure(figsize=(10, 6))
bars = plt.barh(sorted_models, sorted_results, color=['skyblue', 'lightgreen'])
plt.xlabel('Result Percentage')
plt.title('Baseline Model Results')
plt.gca().invert_yaxis() # Invert y-axis to show highest percentage at the top

# Add percentage values on bars
for bar, result in zip(bars, sorted_results):
    plt.text(bar.get_width(), bar.get_y() + bar.get_height()/2, f'{result:.2%}')

plt.show()

```





In [67]:

```


from sklearn.metrics import roc_curve, auc
import numpy as np
import matplotlib.pyplot as plt

# Convert y_pred to numpy array and flatten it
y_pred = np.array(y_pred).flatten()

# Compute ROC curve and ROC area
fpr, tpr, _ = roc_curve(y_true, y_pred)
roc_auc = auc(fpr, tpr)

# Plot ROC curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)'
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()

```



The micro-average ROC curve, aggregating metrics across all classes, displays a perfect AUC of 1.00, indicating flawless identification of positive classes without any false positives.

Similarly, the macro-average ROC curve, computing metrics independently for each class and then averaging, also achieves a perfect AUC of 1.00, reflecting exceptional overall performance.

Examining class-specific ROC curves unveils disparities, notably, Class 0 exhibiting an AUC of 0.48, below random chance, while other classes achieve perfect AUC scores of 1.00. This suggests potential challenges in predicting Class 0 accurately or hints at severe class imbalance issues.

In [39]:

```
# Import necessary Libraries
from sklearn.impute import SimpleImputer

imputer_num = SimpleImputer(strategy='mean')
X_subset_imputed_num = imputer_num.fit_transform(X_subset)
```

In [40]:

```
imputer = SimpleImputer(strategy='mean')
X_train_imputed = imputer.fit_transform(X_train)
X_subset_imputed = imputer.transform(X_subset)
X_test_imputed = imputer.transform(X_test)
```

In [44]:

```
# Assume best_params_brf contains the best parameters found from GridSearchCV
best_params_brf = {'n_estimators': 100, 'max_depth': 10, 'min_samples_split':

# Create the BalancedRandomForestClassifier model with best parameters and up
brf_model = BalancedRandomForestClassifier(**best_params_brf, random_state=42,

# Create the Bagging ensemble using the optimized BalancedRandomForest
bagging_model = BaggingClassifier(estimator=brf_model, n_estimators=10, random

# Fit the Bagging ensemble model to the imputed subset data
bagging_model.fit(X_subset_imputed, y_subset)

# Evaluate the model's performance using the recall metric on the imputed test
y_pred = bagging_model.predict(X_test_imputed)

# Check the unique labels in y_test and y_pred
print("Unique labels in y_test:", np.unique(y_test))
print("Unique labels in y_pred:", np.unique(y_pred))

# Pass the appropriate pos_label based on the unique labels
if 'faulty' in np.unique(y_test) and 'faulty' in np.unique(y_pred):
    pos_label = 'faulty'
else:
    pos_label = 1

print("Recall:", recall_score(y_test, y_pred, pos_label=pos_label))
```

```
Unique labels in y_test: [0 1]
Unique labels in y_pred: [0 1]
Recall: 0.9858206309819213
```

In [47]:

```
from sklearn.impute import SimpleImputer
from sklearn.ensemble import AdaBoostClassifier
from sklearn.metrics import recall_score
from imblearn.ensemble import BalancedRandomForestClassifier

# Instantiate SimpleImputer for numerical features
imputer_num = SimpleImputer(strategy='mean')
X_subset_imputed_num = imputer_num.fit_transform(X_subset)

# Instantiate SimpleImputer
imputer = SimpleImputer(strategy='mean')

# Fit and transform the imputer on the training data
X_train_imputed = imputer.fit_transform(X_train)
```

```

X_train_imputed = imputer.fit_transform(X_train)
X_subset_imputed = imputer.transform(X_subset)
X_test_imputed = imputer.transform(X_test)

# Assume best_params_brf contains the best parameters found from GridSearchCV
best_params_brf = {'n_estimators': 100, 'max_depth': 10, 'min_samples_split':

# Create the BalancedRandomForestClassifier model with best parameters and up
brf_model = BalancedRandomForestClassifier(**best_params_brf, random_state=42,

# Create the AdaBoost classifier using the optimized BalancedRandomForest as t
adaboost_model = AdaBoostClassifier(estimator=brf_model, n_estimators=50, rand

# Fit the AdaBoost model to the imputed subset data
adaboost_model.fit(X_subset_imputed, y_subset)

# Evaluate the model's performance using the recall metric on the imputed test
y_pred = adaboost_model.predict(X_test_imputed)

# Check the unique labels in y_test and y_pred
print("Unique labels in y_test:", np.unique(y_test))
print("Unique labels in y_pred:", np.unique(y_pred))

# Pass the appropriate pos_label based on the unique labels
if 'faulty' in np.unique(y_test) and 'faulty' in np.unique(y_pred):
    pos_label = 'faulty'
else:
    pos_label = 1

print("Recall:", recall_score(y_test, y_pred, pos_label=pos_label))

```

Unique labels in y_test: [0 1]
 Unique labels in y_pred: [0 1]
 Recall: 0.9932647997164126

In [48]:

```

from sklearn.impute import SimpleImputer
import xgboost as xgb
from sklearn.metrics import recall_score

# Instantiate SimpleImputer for numerical features
imputer_num = SimpleImputer(strategy='mean')
X_subset_imputed_num = imputer_num.fit_transform(X_subset)

# Instantiate SimpleImputer
imputer = SimpleImputer(strategy='mean')

# Fit and transform the imputer on the training data
X_train_imputed = imputer.fit_transform(X_train)
X_subset_imputed = imputer.transform(X_subset)
X_test_imputed = imputer.transform(X_test)

# Assume best_params_rf contains the best parameters found from GridSearchCV f
best_params_rf = {'n_estimators': 100, 'max_depth': 10, 'min_samples_split': 2

# Create XGBoost model using parameters similar to the BalancedRandomForest
xgboost_model = xgb.XGBClassifier(n_estimators=best_params_rf['n_estimators'],
                                  max_depth=best_params_rf['max_depth'],
                                  min_child_weight=best_params_rf['min_samples_
                                  subsample=0.8,
                                  colsample_bytree=0.8, # Similar to max_featu
                                  objective='binary:logistic', # Objective for
                                  random_state=42)

# Fit the XGBoost model to the imputed subset data
xgboost_model.fit(X_subset_imputed, y_subset)

```

```

# Evaluate the model's performance using the recall metric on the imputed test
y_pred = xgboost_model.predict(X_test_imputed)

# Check the unique labels in y_test and y_pred
print("Unique labels in y_test:", np.unique(y_test))
print("Unique labels in y_pred:", np.unique(y_pred))

# Pass the appropriate pos_label based on the unique labels
if 'faulty' in np.unique(y_test) and 'faulty' in np.unique(y_pred):
    pos_label = 'faulty'
else:
    pos_label = 1

print("Recall:", recall_score(y_test, y_pred, pos_label=pos_label))

```

```

Unique labels in y_test: [0 1]
Unique labels in y_pred: [0 1]
Recall: 0.9794399149237859

```

various hyperparameters are important for model optimization, such as the number of depth, weight, sub and colsamples. With the tuner initialized for a random search we uncover the best set of hyperparameters, a beacon guiding us towards enhanced model performance. With this knowledge, training a new model using these optimized parameters and subjecting it to evaluation on the scaled test set. Finally, we secure our model, saving it to a designated file path.

In [58]:

```

from sklearn.model_selection import train_test_split
from sklearn.impute import SimpleImputer
from sklearn.ensemble import VotingClassifier
from sklearn.metrics import recall_score

# Split the data into train and test subsets
X_train_subset, X_test_subset, y_train_subset, y_test_subset = train_test_spli

# Impute missing values in the training subset
imputer = SimpleImputer(strategy='mean')
X_train_subset_imputed = imputer.fit_transform(X_train_subset)

# Impute missing values in the test subset
X_test_subset_imputed = imputer.transform(X_test_subset)

# Assuming these models have been optimally configured and trained if necessar
logreg = LogisticRegression(max_iter=1000, solver='saga')
balanced_rf = BalancedRandomForestClassifier(random_state=42, sampling_strateg
bagging = bagging_model # Assuming you have defined bagging_model
adaboost = adaboost_model # Assuming you have defined adaboost_model
xgboost = xgboost_model # Assuming you have defined xgboost_model

# Create the voting classifier
voting_classifier = VotingClassifier(
    estimators=[
        ('logreg', logreg),
        ('balanced_rf', balanced_rf),
        ('bagging', bagging),
        ('adaboost', adaboost),
        ('xgboost', xgboost)
    ],
    voting='soft'
)

# Fit the voting classifier to the imputed training subset
voting_classifier.fit(X_train_subset_imputed, y_train_subset)

```

```
# Make predictions on the imputed test subset
y_pred_subset = voting_classifier.predict(X_test_subset_imputed)

# Evaluate the model's performance on the test subset using recall
print("Recall on subset:", recall_score(y_test_subset, y_pred_subset, pos_label=1))
```

Recall on subset: 0.999291031549096

In [59]:

```
from sklearn.ensemble import StackingClassifier

# Split the data into train and test subsets
X_train_subset, X_test_subset, y_train_subset, y_test_subset = train_test_split(X_train, y_train, test_size=0.2, random_state=42)

# Impute missing values in the training data
imputer = SimpleImputer(strategy='mean')
X_train_imputed = imputer.fit_transform(X_train_subset)

# Impute missing values in the test data
X_test_subset_imputed = imputer.transform(X_test_subset)

# Define a simple meta-learner
meta_learner = LogisticRegression(max_iter=1000)

# Stacking classifier setup
stacking_classifier = StackingClassifier(
    estimators=[
        ('logreg', logreg),
        ('balanced_rf', balanced_rf),
        ('bagging', bagging),
        ('adaboost', adaboost),
        ('xgboost', xgboost)
    ],
    final_estimator=meta_learner
)

# Fit the voting classifier to the imputed training data
stacking_classifier.fit(X_train_subset_imputed, y_train_subset)

# Make predictions on the imputed test data
y_pred_subset = stacking_classifier.predict(X_test_subset_imputed)

# Evaluate the model's performance using recall
print("Recall on subset:", recall_score(y_test_subset, y_pred_subset, pos_label=1))
```

KeyboardInterrupt

Traceback (most recent call last)

Cell In[59], line 29

```
17 stacking_classifier = StackingClassifier(
18     estimators=[
19         ('logreg', logreg),
20     ]
21 )
22     final_estimator=meta_learner
23 )
24 # Fit the voting classifier to the imputed training data
--> 29 stacking_classifier.fit(X_train_subset_imputed, y_train_subset)
30 # Make predictions on the imputed test data
31 y_pred_subset = stacking_classifier.predict(X_test_subset_imputed)
```

File ~\anaconda3\Lib\site-packages\sklearn\ensemble_stacking.py:660, in fit(self, X, y, sample_weight)

```
657     self._label_encoder = [LabelEncoder().fit(yk) for yk in y.T]
658     self.classes_ = [le.classes_ for le in self._label_encoder]
659     y_encoded = np.array(
--> 660         [

```

```

661         self._label_encoder[target_idx].transform(target)
662         for target_idx, target in enumerate(y.T)
663     ]
664     ).T
665 else:
666     self._label_encoder = LabelEncoder().fit(y)

File ~\anaconda3\Lib\site-packages\sklearn\ensemble\_stacking.py:252, in fit(self, X, y, sample_weight)
    241     predictions = [
    242         getattr(estimator, predict_method)(X)
    243         for estimator, predict_method in zip(all_estimators, self.stack_
method_)
    244         if estimator != "drop"
    245     ]
    246 else:
    247     # To train the meta-classifier using the most data as possible, we u
se
    248     # a cross-validation to obtain the output of the stacked estimators.
    249     # To ensure that the data provided to each estimator are the same,
    250     # we need to set the random state of the cv if there is one and we
    251     # need to take a copy.
--> 252     cv = check_cv(self.cv, y=y, classifier=is_classifier(self))
    253     if hasattr(cv, "random_state") and cv.random_state is None:
    254         cv.random_state = np.random.RandomState()

```

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\parallel.py:63, in Parallel.__c
all__(self, iterable)
    58 config = get_config()
    59 iterable_with_config = (
    60     (_with_config(delayed_func, config), args, kwargs)
    61     for delayed_func, args, kwargs in iterable
    62 )
--> 63 return super().__call__(iterable_with_config)

```

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:1088, in Parallel.__call_
(self, iterable)
    1085 if self.dispatch_one_batch(iterator):
    1086     self._iterating = self._original_iterator is not None
-> 1088 while self.dispatch_one_batch(iterator):
    1089     pass
    1091 if pre_dispatch == "all" or n_jobs == 1:
    1092     # The iterable was consumed all at once by the above for loop.
    1093     # No need to wait for async callbacks to trigger to
    1094     # consumption.

```

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:901, in Parallel.dispatch_
one_batch(self, iterator)
    899     return False
    900 else:
--> 901     self._dispatch(tasks)
    902     return True

```

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:819, in Parallel._dispatch
(self, batch)
    817 with self._lock:
    818     job_idx = len(self._jobs)
--> 819     job = self._backend.apply_async(batch, callback=cb)
    820     # A job can complete so quickly than its callback is
    821     # called before we get here, causing self._jobs to
    822     # grow. To ensure correct results ordering, .insert is
    823     # used (rather than .append) in the following line
    824     self._jobs.insert(job_idx, job)

```

```

File ~\anaconda3\Lib\site-packages\joblib\parallel_backends.py:208, in Sequenti
alBackend.apply_async(self, func, callback)
    206 def apply_async(self, func, callback=None):
    207     """Schedule a function in the execution queue"""

```



```

207         """Schedule a func to be run"""
--> 208     result = ImmediateResult(func)
209     if callback:
210         callback(result)

```

File ~\anaconda3\Lib\site-packages\joblib\parallel_backends.py:597, in ImmediateResult.__init__(self, batch)

```

594 def __init__(self, batch):
595     # Don't delay the application, to avoid keeping the input
596     # arguments in memory
--> 597     self.results = batch()

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:288, in BatchedCalls.__call__(self)

```

284 def __call__(self):
285     # Set the default nested backend to self._backend but do not set the
286     # change the default number of processes to -1
287     with parallel_backend(self._backend, n_jobs=self._n_jobs):
--> 288         return [func(*args, **kwargs)
289                 for func, args, kwargs in self.items]

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:288, in <listcomp>(.0)

```

284 def __call__(self):
285     # Set the default nested backend to self._backend but do not set the
286     # change the default number of processes to -1
287     with parallel_backend(self._backend, n_jobs=self._n_jobs):
--> 288         return [func(*args, **kwargs)
289                 for func, args, kwargs in self.items]

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\parallel.py:123, in __call__(self, *args, **kwargs)

```

116 config = getattr(self, "config", None)
117 if config is None:
118     warnings.warn(
119         (
120             "`sklearn.utils.parallel.delayed` should be used with"
121             " `sklearn.utils.parallel.Parallel` to make it possible to"
122             " propagate the scikit-learn configuration of the current th
read to"
--> 123             " the joblib workers."
124         ),
125         UserWarning,
126     )
127     config = {}
128 with config_context(**config):

```

File ~\anaconda3\Lib\site-packages\sklearn\model_selection_validation.py:986, in cross_val_predict(estimator, X, y, groups, cv, n_jobs, verbose, fit_params, pre_dispatch, method)

```

982     scores = scorer(estimator, X_test, y_test, **score_params)
983 except Exception:
984     if isinstance(scorer, _MultimetricScorer):
985         # If `_MultimetricScorer` raises exception, the `error_score`
--> 986         # parameter is equal to "raise".
987         raise
988     else:

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\parallel.py:63, in Parallel.__call__(self, iterable)

```

58 config = get_config()
59 iterable_with_config = (
60     (_with_config(delayed_func, config), args, kwargs)
61     for delayed_func, args, kwargs in iterable
62 )
--> 63 return super().__call__(iterable_with_config)

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:1085, in Parallel.__call__(self, iterable)

```

1076 try:
1077     # Only set self._iterating to True if at least a batch
1078     # was dispatched. In particular this covers the edge
1079     (...)
1080     # was very quick and its callback already dispatched all the
1081     # remaining jobs.
1082     self._iterating = False
-> 1085 if self.dispatch_one_batch(iterator):
1086     self._iterating = self._original_iterator is not None
1088 while self.dispatch_one_batch(iterator):

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:901, in Parallel.dispatch_one_batch(self, iterator)

```

899     return False
900 else:
-> 901     self._dispatch(tasks)
902     return True

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:819, in Parallel._dispatch(self, batch)

```

817 with self._lock:
818     job_idx = len(self._jobs)
-> 819     job = self._backend.apply_async(batch, callback=cb)
820     # A job can complete so quickly that its callback is
821     # called before we get here, causing self._jobs to
822     # grow. To ensure correct results ordering, .insert is
823     # used (rather than .append) in the following line
824     self._jobs.insert(job_idx, job)

```

File ~\anaconda3\Lib\site-packages\joblib\parallel_backends.py:208, in SequentialBackend.apply_async(self, func, callback)

```

206 def apply_async(self, func, callback=None):
207     """Schedule a func to be run"""
-> 208     result = ImmediateResult(func)
209     if callback:
210         callback(result)

```

File ~\anaconda3\Lib\site-packages\joblib\parallel_backends.py:597, in ImmediateResult.__init__(self, batch)

```

594 def __init__(self, batch):
595     # Don't delay the application, to avoid keeping the input
596     # arguments in memory
-> 597     self.results = batch()

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:288, in BatchedCalls.__call__(self)

```

284 def __call__(self):
285     # Set the default nested backend to self._backend but do not set the
286     # change the default number of processes to -1
287     with parallel_backend(self._backend, n_jobs=self._n_jobs):
-> 288         return [func(*args, **kwargs)
289                 for func, args, kwargs in self.items]

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:288, in <listcomp>(.0)

```

284 def __call__(self):
285     # Set the default nested backend to self._backend but do not set the
286     # change the default number of processes to -1
287     with parallel_backend(self._backend, n_jobs=self._n_jobs):
-> 288         return [func(*args, **kwargs)
289                 for func, args, kwargs in self.items]

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\parallel.py:123, in __call__(self, *args, **kwargs)

```

116 config = getattr(self, "config", None)
117 if config is None:
118     warnings.warn(
119         (

```

```

120         "`sklearn.utils.parallel.delayed` should be used with"
121         "`sklearn.utils.parallel.Parallel` to make it possible to"
122         "propagate the scikit-learn configuration of the current th
read to"
--> 123         "the joblib workers."
124     ),
125     UserWarning,
126 )
127     config = {}
128     with config_context(**config):

```

```

File ~\anaconda3\Lib\site-packages\sklearn\model_selection\_validation.py:1068,
in _fit_and_predict(estimator, X, y, train, test, verbose, fit_params, method)
1036         raise ValueError(error_msg % (scores, type(scores), scorer))
1037     return scores
1040 @validate_params(
1041     {
1042         "estimator": [HasMethods(["fit", "predict"])],
1043         "X": ["array-like", "sparse matrix"],
1044         "y": ["array-like", None],
1045         "groups": ["array-like", None],
1046         "cv": ["cv_object"],
1047         "n_jobs": [Integral, None],
1048         "verbose": ["verbose"],
1049         "fit_params": [dict, None],
1050         "params": [dict, None],
1051         "pre_dispatch": [Integral, str, None],
1052         "method": [
1053             StrOptions(
1054                 {
1055                     "predict",
1056                     "predict_proba",
1057                     "predict_log_proba",
1058                     "decision_function",
1059                 }
1060             )
1061     ],
1062 },
1063     prefer_skip_nested_validation=False, # estimator is not validated y
et
1064 )
1065 def cross_val_predict(
1066     estimator,
1067     X,
-> 1068     y=None,
1069     *,
1070     groups=None,
1071     cv=None,
1072     n_jobs=None,
1073     verbose=0,
1074     fit_params=None,
1075     params=None,
1076     pre_dispatch="2*n_jobs",
1077     method="predict",
1078 ):
1079     """Generate cross-validated estimates for each input data point.
1080
1081     The data is split according to the cv parameter. Each sample belongs
1082     (...)
1083
1084     >>> y_pred = cross_val_predict(lasso, X, y, cv=3)
1085     """
1086     params = _check_params_groups_deprecation(fit_params, params, group
s)

```

```

File ~\anaconda3\Lib\site-packages\sklearn\ensemble\_weight_boosting.py:162, in
fit(self, X, y, sample_weight)
159     epsilon = np.finfo(sample_weight.dtype).eps

```

```

161 zero_weight_mask = sample_weight == 0.0
--> 162 for iboost in range(self.n_estimators):
163     # avoid extremely small sample weight, for details see issue #20320
164     sample_weight = np.clip(sample_weight, a_min=epsilon, a_max=None)
165     # do not clip sample weights that were exactly zero originally

```

File ~\anaconda3\Lib\site-packages\sklearn\ensemble_weight_boosting.py:569, in _boost(self, iboost, X, y, sample_weight, random_state)

```

546 def _boost(self, iboost, X, y, sample_weight, random_state):
547     """Implement a single boost.
548
549     Perform a single boost according to the real multi-class SAMME.R
550     algorithm or to the discrete SAMME algorithm and return the updated
551     sample weights.
552
553     Parameters
554     -----
555     iboost : int
556         The index of the current boost iteration.
557
558     X : {array-like, sparse matrix} of shape (n_samples, n_features)
559         The training input samples.
560
561     y : array-like of shape (n_samples,)
562         The target values (class labels).
563
564     sample_weight : array-like of shape (n_samples,)
565         The current sample weights.
566
567     random_state : RandomState instance
568         The RandomState instance used if the base estimator accepts a
--> 569         `random_state` attribute.
570
571     Returns
572     -----
573     sample_weight : array-like of shape (n_samples,) or None
574         The reweighted sample weights.
575         If None then boosting has terminated early.
576
577     estimator_weight : float
578         The weight for the current boost.
579         If None then boosting has terminated early.
580
581     estimator_error : float
582         The classification error for the current boost.
583         If None then boosting has terminated early.
584     """
585     if self.algorithm == "SAMME.R":
586         return self._boost_real(iboost, X, y, sample_weight, random_stat
e)

```

File ~\anaconda3\Lib\site-packages\sklearn\ensemble_weight_boosting.py:578, in _boost_real(self, iboost, X, y, sample_weight, random_state)

```

546 def _boost(self, iboost, X, y, sample_weight, random_state):
547     """Implement a single boost.
548
549     Perform a single boost according to the real multi-class SAMME.R
550     algorithm or to the discrete SAMME algorithm and return the updated
551     sample weights.
552
553     Parameters
554     -----
555     iboost : int
556         The index of the current boost iteration.
557
558     X : {array-like, sparse matrix} of shape (n_samples, n_features)
559         The training input samples.

```

```

560
561     y : array-like of shape (n_samples,)
562         The target values (class labels).
563
564     sample_weight : array-like of shape (n_samples,)
565         The current sample weights.
566
567     random_state : RandomState instance
568         The RandomState instance used if the base estimator accepts a
569         `random_state` attribute.
570
571     Returns
572     -----
573     sample_weight : array-like of shape (n_samples,) or None
574         The reweighted sample weights.
575         If None then boosting has terminated early.
576
577     estimator_weight : float
--> 578         The weight for the current boost.
579         If None then boosting has terminated early.
580
581     estimator_error : float
582         The classification error for the current boost.
583         If None then boosting has terminated early.
584     """
585     if self.algorithm == "SAMME.R":
586         return self._boost_real(iboost, X, y, sample_weight, random_stat
e)

```

File ~\anaconda3\Lib\site-packages\imblearn\utils\fixes.py:85, in _fit_context.<locals>.decorator.<locals>.wrapper(estimator, *args, **kwargs)

```

78     estimator._validate_params()
79     with config_context(
80         skip_parameter_validation=(
81             prefer_skip_nested_validation or global_skip_validation
82         )
83     ):
--> 85     return fit_method(estimator, *args, **kwargs)

```

File ~\anaconda3\Lib\site-packages\imblearn\ensemble_forest.py:676, in BalancedRandomForestClassifier.fit(self, X, y, sample_weight)

```

668     samplers.append(sampler)
669     # Parallel loop: we prefer the threading backend as the Cython code
670     # for fitting the trees is internally releasing the Python GIL
671     # making threading more efficient than multiprocessing in
672     # that case. However, we respect any parallel_backend contexts set
673     # at a higher level, since correctness does not rely on using
674     # threads.
--> 676     samplers_trees = Parallel(
677         n_jobs=self.n_jobs,
678         verbose=self.verbose,
679         prefer="threads",
680     )(
681         delayed(_local_parallel_build_trees)(
682             s,
683             t,
684             self.bootstrap,
685             X,
686             y_encoded,
687             sample_weight,
688             i,
689             len(trees),
690             verbose=self.verbose,
691             class_weight=self.class_weight,
692             n_samples_bootstrap=n_samples_bootstrap,
693             forest=self,
694         )
695     )

```

```

695     for i, (s, t) in enumerate(zip(samplers, trees))
696 )
697 samplers, trees = zip(*samplers_trees)
699 # Collect newly grown trees

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\parallel.py:63, in Parallel.__call__(self, iterable)

```

58 config = get_config()
59 iterable_with_config = (
60     (_with_config(delayed_func, config), args, kwargs)
61     for delayed_func, args, kwargs in iterable
62 )
--> 63 return super().__call__(iterable_with_config)

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:1088, in Parallel.__call__(self, iterable)

```

1085 if self.dispatch_one_batch(iterator):
1086     self.iterating = self._original_iterator is not None
-> 1088 while self.dispatch_one_batch(iterator):
1089     pass
1091 if pre_dispatch == "all" or n_jobs == 1:
1092     # The iterable was consumed all at once by the above for loop.
1093     # No need to wait for async callbacks to trigger to
1094     # consumption.

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:901, in Parallel.dispatch_one_batch(self, iterator)

```

899     return False
900 else:
--> 901     self._dispatch(tasks)
902     return True

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:819, in Parallel._dispatch(self, batch)

```

817 with self._lock:
818     job_idx = len(self._jobs)
--> 819     job = self._backend.apply_async(batch, callback=cb)
820     # A job can complete so quickly that its callback is
821     # called before we get here, causing self._jobs to
822     # grow. To ensure correct results ordering, .insert is
823     # used (rather than .append) in the following line
824     self._jobs.insert(job_idx, job)

```

File ~\anaconda3\Lib\site-packages\joblib\parallel_backends.py:208, in SequentialBackend.apply_async(self, func, callback)

```

206 def apply_async(self, func, callback=None):
207     """Schedule a func to be run"""
--> 208     result = ImmediateResult(func)
209     if callback:
210         callback(result)

```

File ~\anaconda3\Lib\site-packages\joblib\parallel_backends.py:597, in ImmediateResult.__init__(self, batch)

```

594 def __init__(self, batch):
595     # Don't delay the application, to avoid keeping the input
596     # arguments in memory
--> 597     self.results = batch()

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:288, in BatchedCalls.__call__(self)

```

284 def __call__(self):
285     # Set the default nested backend to self._backend but do not set the
286     # change the default number of processes to -1
287     with parallel_backend(self._backend, n_jobs=self._n_jobs):
--> 288         return [func(*args, **kwargs)
289                 for func, args, kwargs in self.items]

```

File ~\anaconda3\Lib\site-packages\joblib\parallel.py:288, in BatchedCalls.__call__(self)

```

116 from ..utils.parallel import _parallel_backend, _parallel_backend_kwargs
117
118 def __call__(self):
119     # Set the default nested backend to self._backend but do not set the
120     # change the default number of processes to -1
121     with parallel_backend(self._backend, n_jobs=self._n_jobs):
122         return [func(*args, **kwargs)
123                 for func, args, kwargs in self.items]

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\parallel.py:123, in __call__(self, *args, **kwargs)

```

116 config = getattr(self, "config", None)
117 if config is None:
118     warnings.warn(
119         (
120             "`sklearn.utils.parallel.delayed` should be used with"
121             "`sklearn.utils.parallel.Parallel` to make it possible to"
122             "propagate the scikit-learn configuration of the current th
read to"
--> 123             "the joblib workers."
124         ),
125         UserWarning,
126     )
127     config = {}
128 with config_context(**config):

```

File ~\anaconda3\Lib\site-packages\imblearn\ensemble_forest.py:65, in _local_parallel_build_trees(sampler, tree, bootstrap, X, y, sample_weight, tree_idx, n_trees, verbose, class_weight, n_samples_bootstrap, forest)

```

47 MAX_INT = np.iinfo(np.int32).max
48 sklearn_version = parse_version(sklearn.__version__)
51 def _local_parallel_build_trees(
52     sampler,
53     tree,
54     bootstrap,
55     X,
56     y,
57     sample_weight,
58     tree_idx,
59     n_trees,
60     verbose=0,
61     class_weight=None,
62     n_samples_bootstrap=None,
63     forest=None,
64     missing_values_in_feature_mask=None,
--> 65 ):
66     # resample before to fit the tree
67     X_resampled, y_resampled = sampler.fit_resample(X, y)
68     if sample_weight is not None:

```

File ~\anaconda3\Lib\site-packages\imblearn\base.py:208, in BaseSampler.fit_resample(self, X, y)

```

187 """Resample the dataset.
188
189 Parameters
190 (...)
191     The corresponding label of `X_resampled`.
192 """
193 self._validate_params()
--> 208 return super().fit_resample(X, y)

```

File ~\anaconda3\Lib\site-packages\imblearn\base.py:112, in SamplerMixin.fit_resample(self, X, y)

```

106 X, y, binarize_y = self._check_X_y(X, y)
107 self.sampling_strategy_ = check_sampling_strategy(
108     self.sampling_strategy, y, self._sampling_type
109 )
--> 112 output = self._fit_resample(X, y)
113 y = (

```

```

115     label_binarize(output[1], classes=np.unique(y)) if binarize_y else o
utput[1]
116 )
118 X_, y_ = arrays_transformer.transform(output[0], y_)

```

File ~\anaconda3\Lib\site-packages\imblearn\under_sampling\prototype_selection_random_under_sampler.py:111, in RandomUnderSampler._fit_resample(self, X, y)

```

107 random_state = check_random_state(self.random_state)
109 idx_under = np.empty((0,), dtype=int)
--> 111 for target_class in np.unique(y):
112     if target_class in self.sampling_strategy_.keys():
113         n_samples = self.sampling_strategy_[target_class]

```

File ~\anaconda3\Lib\site-packages\numpy\lib\arraysetops.py:274, in unique(ar, return_index, return_inverse, return_counts, axis, equal_nan)

```

272 ar = np.asanyarray(ar)
273 if axis is None:
--> 274     ret = _unique1d(ar, return_index, return_inverse, return_counts,
275                     equal_nan=equal_nan)
276     return _unpack_tuple(ret)
278 # axis was specified and not None

```

File ~\anaconda3\Lib\site-packages\numpy\lib\arraysetops.py:336, in _unique1d(ar, return_index, return_inverse, return_counts, equal_nan)

```

334 aux = ar[perm]
335 else:
--> 336     ar.sort()
337     aux = ar
338 mask = np.empty(aux.shape, dtype=np.bool_)

```

KeyboardInterrupt:

In [57]: `print("Recall on subset:", recall_score(y_test_subset, y_pred_subset, pos_label=1))`

Recall on subset: 0.9989365473236441

```

In [69]: import matplotlib.pyplot as plt

# Sample data
models = ['Voting', 'Stacking', 'AdaBoost', 'Bagging', 'XGBoost']
results = [.9992, .9989, .9932, .9858, .9794]

# Sort the results and models in ascending order
sorted_results, sorted_models = zip(*sorted(zip(results, models)))

# Create a multicolor bar graph
plt.figure(figsize=(8, 5))
bars = plt.barh(sorted_models, sorted_results, color=['lightgreen', 'salmon',
plt.xlabel('Result Percentage')
plt.title('Ensemble Model Results')
plt.gca().invert_yaxis() # Invert y-axis to show highest percentage at the top

# Add percentage values on bars
for bar, result in zip(bars, sorted_results):
    plt.text(bar.get_width(), bar.get_y() + bar.get_height()/2, f'{result:.2%}')

plt.show()

```

